

ARMA-GARCH

SHIHENG SHEN

June 6, 2017

1 Temperature Model

1.1 Objective

Build a robust time series model to forecast future temperature's interval.

1.2 Data Source

HadCRUT4 is a gridded dataset of global historical surface temperature anomalies relative to a 1961-1990 reference period. Data are available for each month since January 1850, on a 5 degree grid. http://www.metoffice.gov.uk/hadobs/hadcrut4/data/current/time_series/HadCRUT.4.5.0.0.monthly_ns_avg.txt.

1.3 Data adjustment

From the shown url, grab the monthly global mean of temperature anomalies from 1850-2016, which amounts to a length of 204. Using `decompose()` function in R to get the adjusted value. (See appendix)

Why we didn't use the traditional adjustment method Δ^{12} : I found that `acf(myTS)` shows acf anomaly at lag 24, not at lag 12. And after a Δ^{12} operation, acf at lag 12 became significant. Therefore, we think that using the Δ^{12} is feckless in this case.

1.4 Road Map

We went through a lot of difficulties during the construction of a proper model. Below is the basic idea I've thought about.

- Dealing with Statonarity: `diff()` or time trend(linear, dynamic linear, discrete fourier series)?
- Dealing with stochastic volatility: taking logarithm or exclude some data or both? Stochastic volatility model?
- Choosing model: ARIMA or ARFIMA or ARMA-sGARCH or ARMA-eGARCH or ARFIMA-sGARCH or ARFIMA-eGARCH?

1.5 Starting from *arima* model

First, examine the stationarity of *myTS.adjusted*. Strange enough, it passed the *adf.test* despite that it has a clear trend, as shown in the figure. So we tried to use *auto.arima()* to decide if it is $I(1)$ or $I(0)$. As shown in the code and results, a *arima*(2, 1, 4) model fits *myTS.adjusted* better in both *acf* and *pacf*. The t-value of the coefficients of *auto.arima(residuals)* is smaller for *arima*(2, 1, 4).

The forecast of *arima*(2, 1, 4) for 240 months from 2017.1 is shown in the figure. Within the sample, we compare the forecast and the actual data, only to find that *arima*(2, 1, 4) is far from satisfying.

1.6 Long-Memory Model

Though *acf* shows that some lags are significant, but we think it's only because we took too big a confidential level. Actually the problem is not that great, but we still build a *ARFIMA* model.

We plot the *acf* of the residuals of *ARFIMA* model, but it seems that *ARFIMA* can't effectively explain our data.

1.7 Solving Heterodasticity: Standard GARCH

Looking at the plot of the residuals of *arima*(2, 1, 4). Obviously there exists heterodasticity problem. Using *adf.test*, it supports our doubt. Therefore we have the motive to build a *GARCH* model. We decide to use the package *rugarch*.

A standard GARCH model has the following variance equation:

$$\sigma_t^2 = (w + \sum_{j=1}^m \zeta_j v_{jt}) + \sum_{j=1}^q \alpha_j \varepsilon_{t-j}^2 + \sum_{j=1}^p \beta_j \sigma_{t-j}^2$$

We wrote a function to test different order for the model as shown in appendix. Due to the problem that *rugarch* package only deals with stationary series, we first *diff()* *myTS.adjusted*. After trying a few times from (1, 0), (0, 0, 0) to (5, 5), (4, 0, 4), *GARCH*(1, 1), *ARIMA*(2, 0, 3) is the most satisfying model. Here we realize that using GARCH model, the order of ARIMA might changes. Subtract the standardized(w.r.t. the variance model) residuals z , which is $z = \frac{\text{residuals}(fit)}{\text{sigma}(fit)}$. Plot z , we'll find z seems to have a far smoother variance. Just for fun, plot a normal sample series with the same parameters. We'll find in the two figures it's hard to distinguish between the two. Looking at the *LM - Test* in the summary of *fit1*, the heterodasticity is wiped out at lower lags. More graphical diagnostics is available in the figure appendix. From these diagnostics we find sGARCH is to some extent a proper model. Since we difference the series at first, now forecast becomes more complex. We have to calculate the variance by adding all the terms' variance.

2 Appendix

2.1 R Code

2.1.1 Data Adjustment

```
library(curl)
tmpf <- tempfile()
curl_download(url, tmpf)
gtemp <- read.table(tmpf)[, 1:2]
temp = gtemp$V2[1:2004]
library(TSA)
myTS = ts(as.numeric(temp), start = c(1850, 1), frequency = 12)
myTS.additive = decompose(myTS)
myTS.adjusted = as.numeric(myTS.additive$x - myTS.additive$seasonal)
```

2.1.2 arima models

```
> tseries::adf.test(myTS.adjusted)
```

Augmented Dickey-Fuller Test

```
data: myTS.adjusted
Dickey-Fuller = -5.1646, Lag order = 12, p-value = 0.01
alternative hypothesis: stationary
```

```
> library(forecast)
> auto.arima(myTS.adjusted)
Series: myTS.adjusted
ARIMA(2,1,4)(2,0,1)[12] with drift
```

Coefficients:

	ar1	ar2	ma1	ma2	ma3	ma4	sar1	sar2	sma1
	0.5041	0.3585	-1.0458	-0.1600	0.1362	0.0779	0.7619	0.0779	-0.7884
s.e.	0.2432	0.2143	0.2424	0.3482	0.1082	0.0257	0.0926	0.0248	0.0910
	drift								
	5e-04								
s.e.	2e-04								

```
sigma^2 estimated as 0.01428: log likelihood=1417.03
AIC=-2812.05 AICc=-2811.92 BIC=-2750.43
```

```
> arima1 = Arima(myTS.adjusted, c(2, 0, 1))
> auto.arima(arima1$residuals)
Series: arima1$residuals
ARIMA(5,1,0)(2,0,1)[12]
```

Coefficients:

	ar1	ar2	ar3	ar4	ar5	sar1	sar2	sma1
	-0.8556	-0.6256	-0.4849	-0.3136	-0.1475	0.7647	0.0837	-0.7907
s.e.	0.0224	0.0287	0.0299	0.0286	0.0222	0.0846	0.0248	0.0831

```
sigma^2 estimated as 0.01705: log likelihood=1238.65
AIC=-2459.29 AICc=-2459.2 BIC=-2408.87
```

```
> arima2 = Arima(myTS.adjusted, c(2, 1, 4))
> auto.arima(arima2$residuals)
Series: arima2$residuals
ARIMA(1,0,1)(2,0,1)[12] with non-zero mean
```

```

Coefficients:
      ar1      ma1      sar1      sar2      sma1      mean
    0.5645 -0.5685  0.7389  0.0745 -0.7672  0.0062
s.e.  2.5541  2.5131  0.1253  0.0244  0.1242  0.0033

```

```

sigma^2 estimated as 0.01427:  log likelihood=1417.4
AIC=-2820.81  AICc=-2820.75  BIC=-2781.59

```

```

# looking at acf
# Choose arima(2, 1, 4)
> tseries::adf.test(arima2$residuals)

```

Augmented Dickey-Fuller Test

```

data:  arima2$residuals
Dickey-Fuller = -11.777, Lag order = 12,
p-value = 0.01
alternative hypothesis: stationary

```

```

> arima2
Series: myTS.adjusted
ARIMA(2,1,4)

```

```

Coefficients:
      ar1      ar2      ma1      ma2      ma3      ma4
    0.5129  0.3271 -1.0438 -0.1331  0.1166  0.0748
s.e.  0.2753  0.2357  0.2748  0.3833  0.1117  0.0265

```

```

sigma^2 estimated as 0.01443:  log likelihood=1405.3
AIC=-2796.6  AICc=-2796.55  BIC=-2757.38

```

```

#forecast for future
> plot(forecast.Arima(arima2, h = 240))

# forecast within the sample and comparision
> sarima = Arima(myTS.adjusted[1:1800], c(2, 1, 4))
> plot(forecast.Arima(sarima, h = 203))
> lines(myTS.adjusted)

```

2.1.3 ARFIMA Model

```

> lmodel = arfima(myTS.adjusted)
> summary(lmodel)

```

```

Call:
  arfima(y = myTS.adjusted)

```

```

*** Warning during (fdcov) fit: unable to compute correlation matrix; maybe change 'h'

```

```

Coefficients:
      Estimate
d           0.445
ar.ar1      0.249
ar.ar2      0.612
ma.ma1      0.223
ma.ma2      0.549
sigma[eps] = 0.1200103

```

```
[d.tol = 0.0001221, M = 100, h = 1.481e-05]
Log likelihood: 1404 ==> AIC = -2796.513 [6 deg.freedom]

> acf(lmodel$residuals)
```

2.1.4 sGARCH Model

```
> arch.test(arma2$residuals)
ARCH heteroscedasticity test for residuals
alternative: heteroscedastic
```

Portmanteau-Q test:

	order	PQ	p.value
[1,]	4	99.4	0
[2,]	8	116.9	0
[3,]	12	411.5	0
[4,]	16	475.6	0
[5,]	20	495.9	0
[6,]	24	755.8	0

Lagrange-Multiplier test:

	order	LM	p.value
[1,]	4	1418	0
[2,]	8	697	0
[3,]	12	418	0
[4,]	16	212	0
[5,]	20	168	0
[6,]	24	137	0

```
library(rugarch)
my_sGARCH_test <- function(p, q, m, n, ts.data = res)
{
  # I use include.mean = FALSE after trying TRUE
  # to find out insignificance
  myspec=ugarchspec(variance.model = list(model = "sGARCH", garchOrder = c(p, q)),
    mean.model = list(armaOrder = c(m, n), include.mean = FALSE),
    distribution.model = "normal")
  myfit=ugarchfit(myspec,data=ts.data, solver="solnp")
  return(myfit)
}
```

```
> dtemp = diff(myTS.adjusted)
> fit1 = my_sGARCH_test(1, 1, 2, 3, dtemp)
```

```
> fit1
```

```
*-----*
*          GARCH Model Fit          *
*-----*
```

Conditional Variance Dynamics

```
-----
GARCH Model : sGARCH(1,1)
Mean Model : ARFIMA(2,0,3)
Distribution : norm
```

Optimal Parameters

	Estimate	Std. Error	t value	Pr(> t)
ar1	-0.090935	0.015226	-5.9724	0.000000
ar2	0.754333	0.015069	50.0583	0.000000
ma1	-0.391906	0.007313	-53.5917	0.000000
ma2	-0.863257	0.000130	-6647.5909	0.000000
ma3	0.295634	0.007800	37.9010	0.000000
omega	0.000082	0.000034	2.3956	0.016595
alpha1	0.023026	0.003614	6.3706	0.000000
beta1	0.970594	0.004705	206.2925	0.000000

Robust Standard Errors:

	Estimate	Std. Error	t value	Pr(> t)
ar1	-0.090935	0.017031	-5.3393	0.000000
ar2	0.754333	0.017312	43.5726	0.000000
ma1	-0.391906	0.002647	-148.0378	0.000000
ma2	-0.863257	0.000144	-6000.9074	0.000000
ma3	0.295634	0.002903	101.8443	0.000000
omega	0.000082	0.000036	2.2838	0.022382
alpha1	0.023026	0.003636	6.3329	0.000000
beta1	0.970594	0.003734	259.9413	0.000000

LogLikelihood : 1512.336

Information Criteria

Akaike	-1.5021
Bayes	-1.4797
Shibata	-1.5021
Hannan-Quinn	-1.4939

Weighted Ljung-Box Test on Standardized Residuals

	statistic	p-value
Lag[1]	0.3742	0.5407
Lag[2*(p+q)+(p+q)-1] [14]	4.5219	1.0000
Lag[4*(p+q)+(p+q)-1] [24]	13.5547	0.3236
d.o.f=5		
H0 : No serial correlation		

Weighted Ljung-Box Test on Standardized Squared Residuals

	statistic	p-value
Lag[1]	31.06	2.506e-08
Lag[2*(p+q)+(p+q)-1] [5]	37.84	1.904e-10
Lag[4*(p+q)+(p+q)-1] [9]	51.28	4.433e-13
d.o.f=2		

Weighted ARCH LM Tests

	Statistic	Shape	Scale	P-Value
ARCH Lag[3]	0.004507	0.500	2.000	9.465e-01
ARCH Lag[5]	11.085543	1.440	1.667	3.565e-03
ARCH Lag[7]	20.592184	2.315	1.543	4.509e-05

Nyblom stability test

Joint Statistic: 2.6294

Individual Statistics:

ar1 0.25019
ar2 0.61592
ma1 0.19513
ma2 0.24162
ma3 0.07259
omega 0.10095
alpha1 0.42035
beta1 0.20019

Asymptotic Critical Values (10% 5% 1%)

Joint Statistic: 1.89 2.11 2.59

Individual Statistic: 0.35 0.47 0.75

Sign Bias Test

 t-value prob sig
Sign Bias 0.3363 7.367e-01
Negative Sign Bias 3.6815 2.380e-04 ***
Positive Sign Bias 4.4421 9.396e-06 ***
Joint Effect 33.2983 2.786e-07 ***

Adjusted Pearson Goodness-of-Fit Test:

group statistic p-value(g-1)
1 20 57.32 1.019e-05
2 30 63.59 2.167e-04
3 40 85.01 2.883e-05
4 50 102.77 1.103e-05

Elapsed time : 0.365526

```
> z = residuals(fit1) / sigma(fit1)
> plot.ts(z)
> mean(z)
[1] 0.03181525
> var(z)
[1] 1.013866
> length(z)
[1] 2003
> plot.ts(rnorm(2003, 0.03181525, 1.013866))

# forecast
> fore1 = ugarchforecast(fit1, n.ahead = 24)
> fore.diff = as.numeric(fore1@forecast$seriesFor)
> fore.sigma = as.numeric(fore1@forecast$sigmaFor)
> ts.predict = temp[length(temp)] + cumsum(fore.diff)
> ts.predict = ts.predict + myTS.additive$figure
> ts.sigma = sqrt(cumsum(fore.sigma^2))
> tsup.sigma = ts.predict + ts.sigma
> tsdown.sigma = ts.predict - ts.sigma
> plot(1:24, ts.predict, ylim=c(0,1.5), type = 'l', col = 'blue',
      xlab = "months", ylab = "temperature predict")
```

```
> lines(1:24, tsup.sigma, type = 'l', col = 'red')
> lines(1:24, ttdown.sigma, type = 'l', col = 'red')
```

2.2 Figures

Figure 1: HadCRUT4 Data Global Mean Time Series

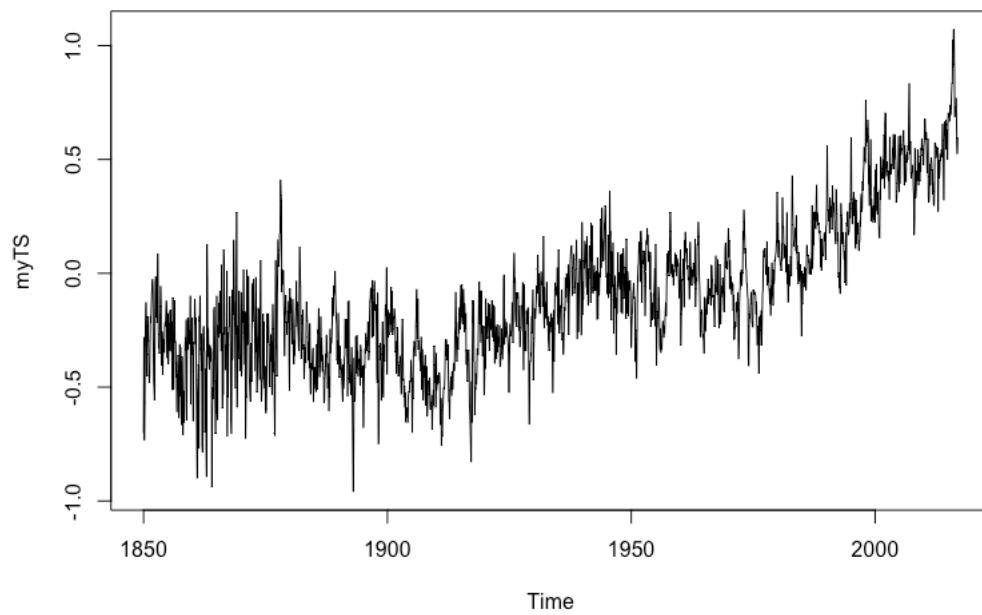


Figure 2: Plot myTS.adjusted




Figure 3: Selecting arma order: acf of residuals of arma(2, 0, 1) & (2, 1, 4)

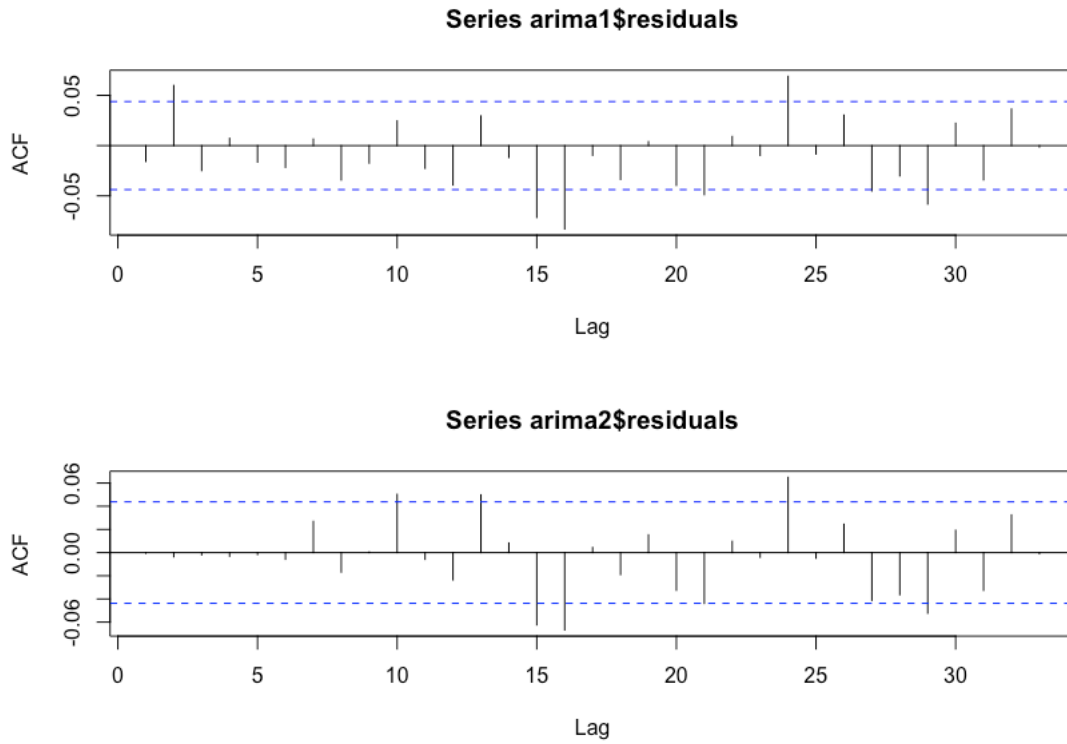


Figure 4: Selecting arma order: pacf of residuals of arma(2, 0, 1) & (2, 1, 4)

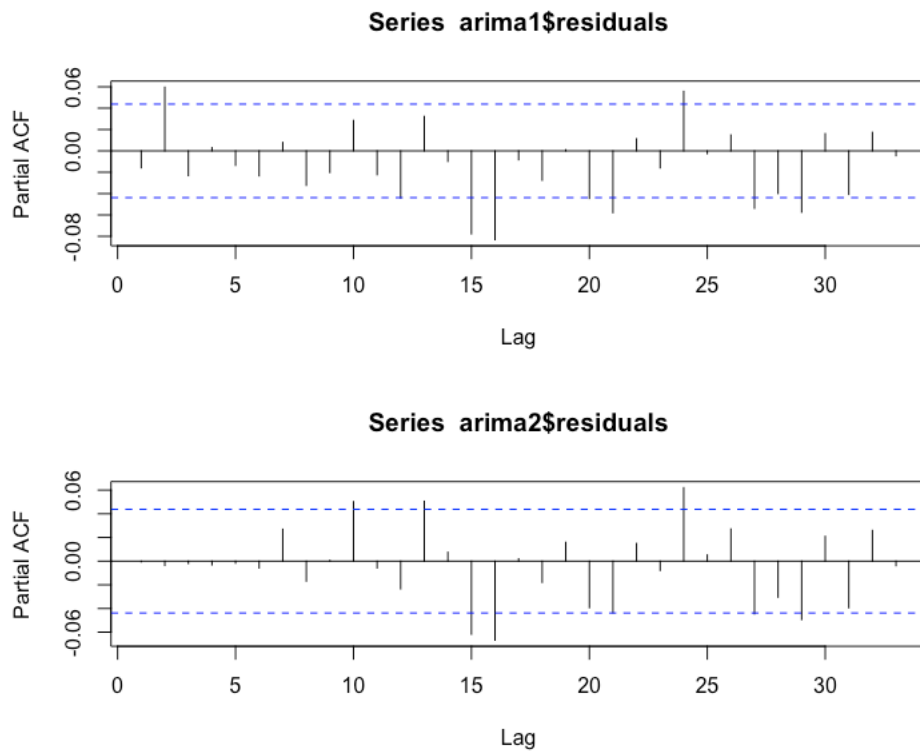


Figure 5: Forecast next 20 years(240 months) using arima(2, 1, 4)




Figure 6: Forecast within the sample compared with the actual data using arima(2, 1, 4)




Figure 7: Plot arima2's residuals




Figure 8: acf of residuals of ARFIMA

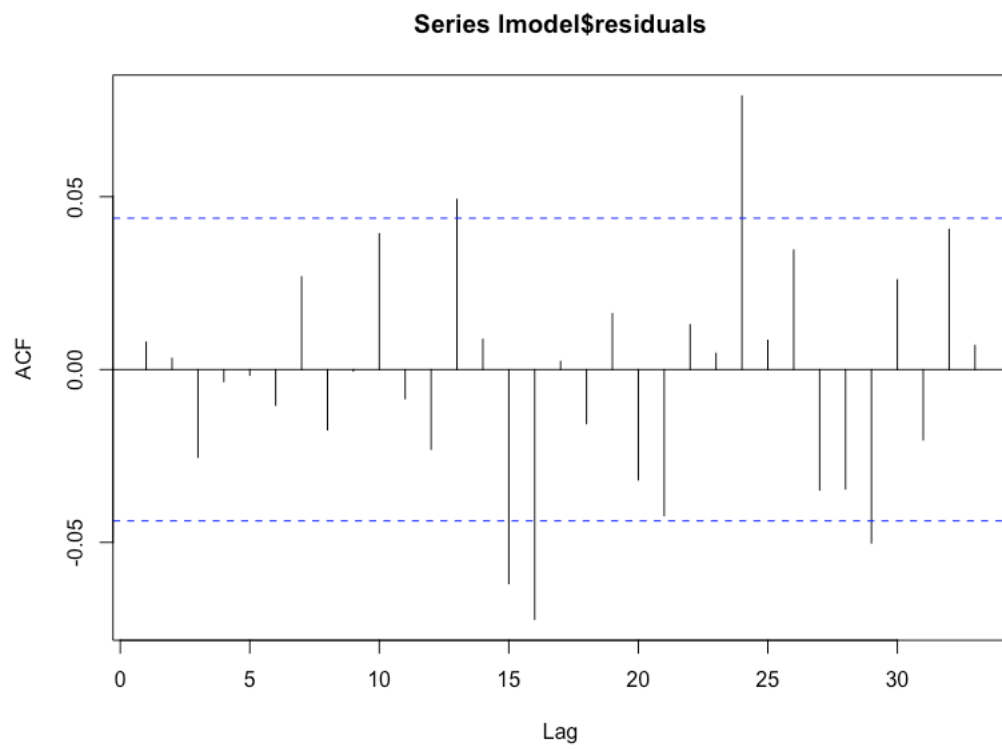


Figure 9: z




Figure 10: simulation of using `rnorm`




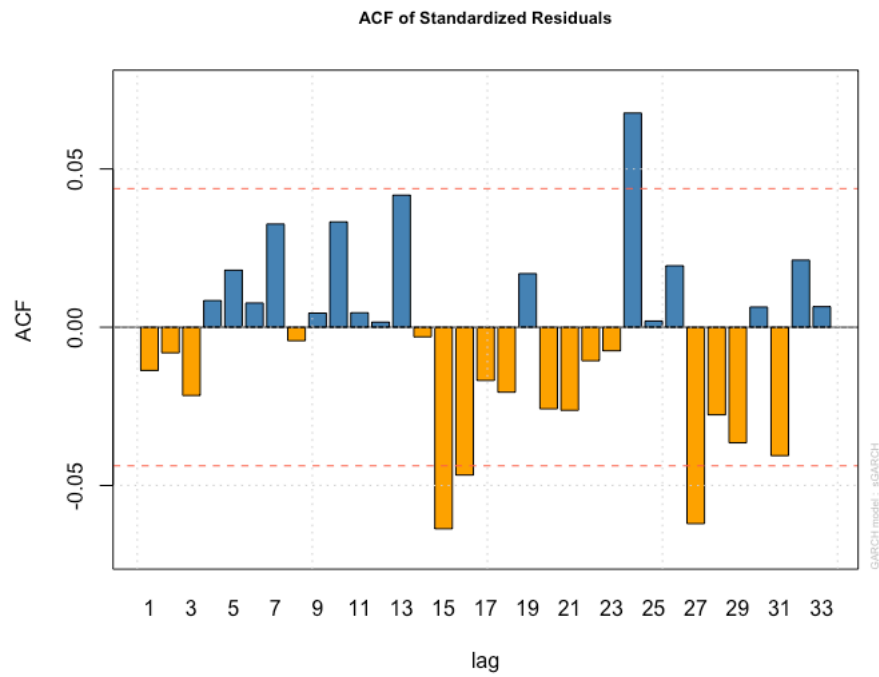
Figure 11: $\text{acf}(z)$ 

Figure 12: Empirical Density of Standardized Residuals

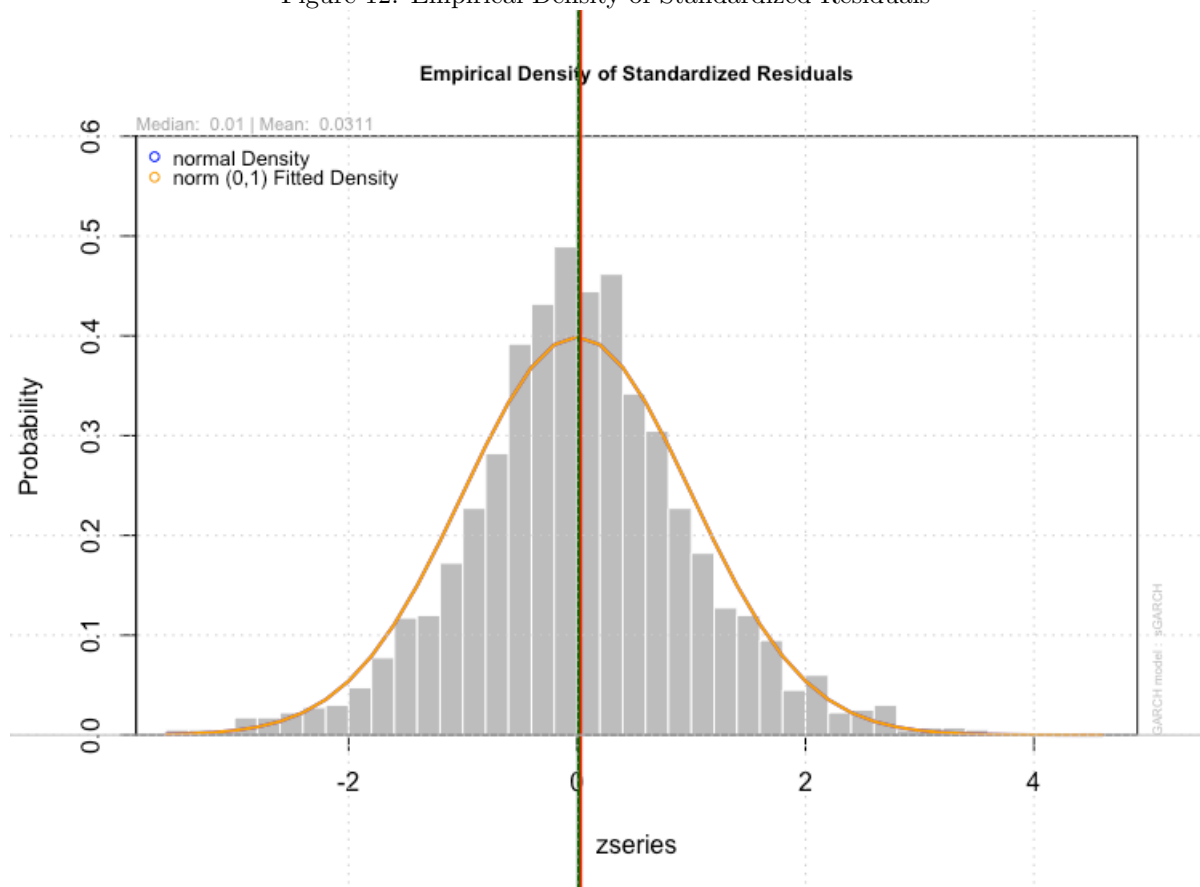


Figure 13: qqplot of z





Figure 14: forecast for sGARCH Model



2.3 Others

2.3.1 Details about *decompose()* used in seasonal adjustment

Type 'decompose' in R console, we can see the source code of this function. The process of 'type = additive' is listed below:

- the argument passed into *decompose()* is a 'ts' object
- denote the argument *ts(x, frequency = f)*. Create a filter using: $\text{filter} = c(0.5, \text{rep}(1, f - 1), 0.5)/f$.
- $\text{trend} = \text{filter}(x, \text{filter})$
- $\text{season} = x - \text{trend}$, then compute f means of season with interval length f , the f means denoted by figure. Adjusting figure $\text{figure} = \text{figure} - \text{mean}(\text{figure})$
- seasonal is just $\text{length}(x)/f$ times repetition of figure.
- $\text{random} = x - \text{seasonal} - \text{trend}$