

ARMA-GARCH

SHIHENG SHEN

June 10, 2017

1 Temperature Model

1.1 Objective

Build a robust time series model to forecast future temperature's interval.

1.2 Data Source

HadCRUT4 is a gridded dataset of global historical surface temperature anomalies relative to a 1961-1990 reference period. Data are available for each month since January 1850, on a 5 degree grid. http://www.metoffice.gov.uk/hadobs/hadcrut4/data/current/time_series/HadCRUT.4.5.0.0.monthly_ns_avg.txt.

1.3 Data adjustment

From the shown url, grab the monthly global mean of temperature anomalies from 1850-2016, which amounts to a length of 204. Using `decompose()` function in R to get the adjusted value. (See appendix)

Why we didn't use the traditional adjustment method Δ^{12} : I found that `acf(myTS)` shows acf anomaly at lag 24, not at lag 12. And after a Δ^{12} operation, acf at lag 12 became significant. Therefore, we think that using the Δ^{12} is feckless in this case.

1.4 Starting from *arima* model

First, examine the stationarity of *myTS.adjusted*. Strange enough, it passed the *adf.test* despite that it has a clear trend, as shown in the figure. So we tried to use `auto.arima()` to decide if it is $I(1)$ or $I(0)$. As shown in the code and results, a *arima*(2, 1, 4) model fits *myTS.adjusted* better in both *acf* and *pacf*. The t-value of the coefficients of `auto.arima(residuals)` is smaller for *arima*(2, 1, 4).

The forecast of *arima*(2, 1, 4) for 240 months from 2017.1 is shown in the figure. Within the sample, we compare the forecast and the actual data, only to find that *arima*(2, 1, 4) is far from satisfying.

1.5 Long-Memory Model

Though *acf* shows that some lags are significant, but we think it's only because we took too big a confidence level. Actually the problem is not that great, but we still build a *ARFIMA* model.

We plot the *acf* of the residuals of *ARFIMA* model, but it seems that *ARFIMA* can't effectively explain our data.

1.6 More remark on the long-memory effect

We find that *acf* of residuals always display significance at lag 24, which cannot be eliminated by long-memory model. Also, as mentioned in the 'Data adjustment part', seasonal adjustment with *period* = 12(12 months) cannot solve the problem as well. We've also tried to adjust the effect using 24 lags, but to find it ineffective as well. It should also be noticed that using `auto.arima()` with seasonal set to TRUE is also bootless, because there is only lag 24 significant, and these 'auto' functions only accept a logical seasonal value, therefore I cannot assign the fixed orders. At last, I find that I can manually find a somehow acceptable method to solve the 24 lag. (but there's another lag becoming significant, at lag 26 or something) The code and figures are in the appendix

Since I do not know how to combine this method with GARCH model, I abandon it when dealing with heterodasticity problem.

1.7 Solving Heterodasticity: Standard GARCH

Looking at the plot of the residuals of *arima*(2, 1, 4). Obviously there exists heterodasticity problem. Using *adf.test*, it supports our doubt. Therefore we have the motive to build a *GARCH* model. We decide to use the package *rugarch*.

A standard GARCH model has the following variance equation:

$$\sigma_t^2 = (w + \sum_{j=1}^m \zeta_j v_{jt}) + \sum_{j=1}^q \alpha_j \varepsilon_{t-j}^2 + \sum_{j=1}^p \beta_j \sigma_{t-j}^2$$

We wrote a function to test different order for the model as shown in appendix. Due to the problem that *rugarch* package only deals with stationary series, we first *diff()* *myTS.adjusted*. After trying a few times from (1, 0), (0, 0, 0) to (5, 5), (4, 0, 4), *GARCH*(1, 1), *ARIMA*(2, 0, 3) is the most satisfying model. Here we realize that using GARCH model, the order of ARIMA might changes. Subtract the standardized(w.r.t. the variance model) residuals z , which is $z = \frac{residuals(fit)}{sigma(fit)}$. Plot z , we'll find z seems to have a far smoother variance. Just for fun, plot a normal sample series with the same parameters. We'll find in the two figures it's hard to distinguish between the two. Looking at the *LM - Test* in the summary of *fit1*, the heterodasticity is wiped out at lower lags. More graphical diagnostics is available in the figure appendix. From these diagnostics we find sGARCH is to some extent a proper model. Since we difference the series at first, now forecast becomes more complex. We have to calculate the variance by adding all the terms' variance.

2 Appendix

2.1 R Code

2.1.1 Data Adjustment

```
library(curl)
tmpf <- tempfile()
curl_download(url, tmpf)
gtemp <- read.table(tmpf)[, 1:2]
temp = gtemp$V2[1:2004]
library(TSA)
myTS = ts(as.numeric(temp), start = c(1850, 1), frequency = 12)
myTS.additive = decompose(myTS)
myTS.adjusted = as.numeric(myTS.additive$x - myTS.additive$seasonal)
```

2.1.2 arima models

```
> tseries::adf.test(myTS.adjusted)
```

Augmented Dickey-Fuller Test

```
data: myTS.adjusted
Dickey-Fuller = -5.1646, Lag order = 12, p-value = 0.01
alternative hypothesis: stationary
```

```
> library(forecast)
> auto.arima(myTS.adjusted)
Series: myTS.adjusted
ARIMA(2,1,4)(2,0,1)[12] with drift
```

Coefficients:

	ar1	ar2	ma1	ma2	ma3	ma4	sar1	sar2	sma1
	0.5041	0.3585	-1.0458	-0.1600	0.1362	0.0779	0.7619	0.0779	-0.7884
s.e.	0.2432	0.2143	0.2424	0.3482	0.1082	0.0257	0.0926	0.0248	0.0910
	drift								
	5e-04								
s.e.	2e-04								

```
sigma^2 estimated as 0.01428: log likelihood=1417.03
AIC=-2812.05 AICc=-2811.92 BIC=-2750.43
```

```
> arima1 = Arima(myTS.adjusted, c(2, 0, 1))
> auto.arima(arima1$residuals)
Series: arima1$residuals
ARIMA(5,1,0)(2,0,1)[12]
```

Coefficients:

	ar1	ar2	ar3	ar4	ar5	sar1	sar2	sma1
	-0.8556	-0.6256	-0.4849	-0.3136	-0.1475	0.7647	0.0837	-0.7907
s.e.	0.0224	0.0287	0.0299	0.0286	0.0222	0.0846	0.0248	0.0831

```
sigma^2 estimated as 0.01705: log likelihood=1238.65
AIC=-2459.29 AICc=-2459.2 BIC=-2408.87
```

```
> arima2 = Arima(myTS.adjusted, c(2, 1, 4))
> auto.arima(arima2$residuals)
Series: arima2$residuals
ARIMA(1,0,1)(2,0,1)[12] with non-zero mean
```

```

Coefficients:
      ar1      ma1      sar1      sar2      sma1      mean
    0.5645 -0.5685  0.7389  0.0745 -0.7672  0.0062
s.e.  2.5541  2.5131  0.1253  0.0244  0.1242  0.0033

```

```

sigma^2 estimated as 0.01427:  log likelihood=1417.4
AIC=-2820.81  AICc=-2820.75  BIC=-2781.59

```

```

# looking at acf
# Choose arima(2, 1, 4)
> tseries::adf.test(arima2$residuals)

```

Augmented Dickey-Fuller Test

```

data:  arima2$residuals
Dickey-Fuller = -11.777, Lag order = 12,
p-value = 0.01
alternative hypothesis: stationary

```

```

> arima2
Series: myTS.adjusted
ARIMA(2,1,4)

```

```

Coefficients:
      ar1      ar2      ma1      ma2      ma3      ma4
    0.5129  0.3271 -1.0438 -0.1331  0.1166  0.0748
s.e.  0.2753  0.2357  0.2748  0.3833  0.1117  0.0265

```

```

sigma^2 estimated as 0.01443:  log likelihood=1405.3
AIC=-2796.6  AICc=-2796.55  BIC=-2757.38

```

```

#forecast for future
> plot(forecast.Arima(arima2, h = 240))

# forecast within the sample and comparision
> sarima = Arima(myTS.adjusted[1:1800], c(2, 1, 4))
> plot(forecast.Arima(sarima, h = 203))
> lines(myTS.adjusted)

```

2.1.3 ARFIMA Model

```

> lmodel = arfima(myTS.adjusted)
> summary(lmodel)

```

```

Call:
  arfima(y = myTS.adjusted)

```

```

*** Warning during (fdcov) fit: unable to compute correlation matrix; maybe change 'h'

```

```

Coefficients:
      Estimate
d           0.445
ar.ar1      0.249
ar.ar2      0.612
ma.ma1      0.223
ma.ma2      0.549
sigma[eps] = 0.1200103

```

```
[d.tol = 0.0001221, M = 100, h = 1.481e-05]
Log likelihood: 1404 ==> AIC = -2796.513 [6 deg.freedom]
```

```
> acf(lmodel$residuals)
```

2.1.4 Dealing with the 24-lag anomaly

```
> gtemp = as.numeric(gtemp$V2)
Error in gtemp$V2 : $ operator is invalid for atomic vectors
> mytemp = gtemp[600:1980]
> plot.ts(mytemp)
> acf(mytemp)
> pacf(mytemp)
> m1 = auto.arima(mytemp, seasonal = TRUE)
> m1
Series: mytemp
ARIMA(3,1,1) with drift

Coefficients:
          ar1      ar2      ar3      ma1  drift
      0.4766  0.2072  0.0823 -0.9809  7e-04
s.e.  0.0276  0.0296  0.0274  0.0062  2e-04

sigma^2 estimated as 0.01009: log likelihood=1215.01
AIC=-2418.01 AICc=-2417.95 BIC=-2386.64
> m2 = arima(mytemp, order = c(3, 1, 1), seasonal = list(order = c(1, 0, 1), period = 24))
> m2
```

```
Call:
arima(x = mytemp, order = c(3, 1, 1), seasonal = list(order = c(1, 0, 1), period = 24))
```

```
Coefficients:
          ar1      ar2      ar3      ma1      sar1      sma1
      0.4728  0.2125  0.0850 -0.9783  0.9232 -0.8713
s.e.  0.0273  0.0294  0.0273  0.0063  0.0397  0.0518
```

```
sigma^2 estimated as 0.009917: log likelihood = 1223.8, aic = -2435.6
> m3 = arfima(mytemp)
> m3
```

```
Call:
arfima(y = mytemp)
```

```
Coefficients:
          d      ar.ar1      ar.ar2      ma.ma1
0.49203791 0.70817900 0.08482127 0.68963222
sigma[eps] = 0.1010365
a list with components:
 [1] "log.likelihood" "n"          "msg"
 [4] "d"              "ar"          "ma"
 [7] "covariance.dpq" "fnormMin"    "sigma"
[10] "stderror.dpq"   "correlation.dpq" "h"
[13] "d.tol"          "M"           "hessian.dpq"
[16] "length.w"       "call"        "residuals"
[19] "x"              "fitted"      "series"
> res3 = residuals(m3)
> res4 = residuals(arima(res3, seasonal = list(order = c(1, 0, 1), period = 24)))
> acf(res4)
```

```
# try to use pure seasonal adjustment with period=24
> atemp = diff(mytemp, lag = 24)
> m5 = auto.arima(atemp)
> res5 = residuals(m5)
> acf(res5)
```

2.1.5 sGARCH Model

```
> arch.test(arma2$residuals)
ARCH heteroscedasticity test for residuals
alternative: heteroscedastic
```

Portmanteau-Q test:

	order	PQ	p.value
[1,]	4	99.4	0
[2,]	8	116.9	0
[3,]	12	411.5	0
[4,]	16	475.6	0
[5,]	20	495.9	0
[6,]	24	755.8	0

Lagrange-Multiplier test:

	order	LM	p.value
[1,]	4	1418	0
[2,]	8	697	0
[3,]	12	418	0
[4,]	16	212	0
[5,]	20	168	0
[6,]	24	137	0

```
library(rugarch)
my_sGARCH_test <- function(p, q, m, n, ts.data = res)
{
  # I use include.mean = FALSE after trying TRUE
  # to find out insignificance
  myspec=ugarchspec(variance.model = list(model = "sGARCH", garchOrder = c(p, q)),
    mean.model = list(armaOrder = c(m, n), include.mean = FALSE),
    distribution.model = "normal")
  myfit=ugarchfit(myspec,data=ts.data, solver="solnp")
  return(myfit)
}
```

```
> dtemp = diff(myTS.adjusted)
> fit1 = my_sGARCH_test(1, 1, 2, 3, dtemp)
```

```
> fit1
```

```
*-----*
*          GARCH Model Fit          *
*-----*
```

Conditional Variance Dynamics

```
-----
GARCH Model : sGARCH(1,1)
Mean Model : ARFIMA(2,0,3)
Distribution : norm
```

Optimal Parameters

```
-----
```

	Estimate	Std. Error	t value	Pr(> t)
ar1	-0.090935	0.015226	-5.9724	0.000000
ar2	0.754333	0.015069	50.0583	0.000000
ma1	-0.391906	0.007313	-53.5917	0.000000
ma2	-0.863257	0.000130	-6647.5909	0.000000
ma3	0.295634	0.007800	37.9010	0.000000
omega	0.000082	0.000034	2.3956	0.016595
alpha1	0.023026	0.003614	6.3706	0.000000
beta1	0.970594	0.004705	206.2925	0.000000

Robust Standard Errors:

	Estimate	Std. Error	t value	Pr(> t)
ar1	-0.090935	0.017031	-5.3393	0.000000
ar2	0.754333	0.017312	43.5726	0.000000
ma1	-0.391906	0.002647	-148.0378	0.000000
ma2	-0.863257	0.000144	-6000.9074	0.000000
ma3	0.295634	0.002903	101.8443	0.000000
omega	0.000082	0.000036	2.2838	0.022382
alpha1	0.023026	0.003636	6.3329	0.000000
beta1	0.970594	0.003734	259.9413	0.000000

LogLikelihood : 1512.336

Information Criteria

Akaike	-1.5021
Bayes	-1.4797
Shibata	-1.5021
Hannan-Quinn	-1.4939

Weighted Ljung-Box Test on Standardized Residuals

	statistic	p-value
Lag[1]	0.3742	0.5407
Lag[2*(p+q)+(p+q)-1] [14]	4.5219	1.0000
Lag[4*(p+q)+(p+q)-1] [24]	13.5547	0.3236

d.o.f=5
H0 : No serial correlation

Weighted Ljung-Box Test on Standardized Squared Residuals

	statistic	p-value
Lag[1]	31.06	2.506e-08
Lag[2*(p+q)+(p+q)-1] [5]	37.84	1.904e-10
Lag[4*(p+q)+(p+q)-1] [9]	51.28	4.433e-13

d.o.f=2

Weighted ARCH LM Tests

	Statistic	Shape	Scale	P-Value
ARCH Lag[3]	0.004507	0.500	2.000	9.465e-01
ARCH Lag[5]	11.085543	1.440	1.667	3.565e-03
ARCH Lag[7]	20.592184	2.315	1.543	4.509e-05

Nyblom stability test

Joint Statistic: 2.6294

Individual Statistics:

```
ar1    0.25019
ar2    0.61592
ma1    0.19513
ma2    0.24162
ma3    0.07259
omega  0.10095
alpha1 0.42035
beta1  0.20019
```

Asymptotic Critical Values (10% 5% 1%)

```
Joint Statistic:      1.89 2.11 2.59
Individual Statistic: 0.35 0.47 0.75
```

Sign Bias Test

```
-----
                t-value      prob sig
Sign Bias      0.3363 7.367e-01
Negative Sign Bias 3.6815 2.380e-04 ***
Positive Sign Bias 4.4421 9.396e-06 ***
Joint Effect    33.2983 2.786e-07 ***
```

Adjusted Pearson Goodness-of-Fit Test:

```
-----
group statistic p-value(g-1)
1    20      57.32   1.019e-05
2    30      63.59   2.167e-04
3    40      85.01   2.883e-05
4    50     102.77   1.103e-05
```

Elapsed time : 0.365526

```
> z = residuals(fit1) / sigma(fit1)
> plot.ts(z)
> mean(z)
[1] 0.03181525
> var(z)
[1] 1.013866
> length(z)
[1] 2003
> plot.ts(rnorm(2003, 0.03181525, 1.013866))

# forecast
> fore1 = ugarchforecast(fit1, n.ahead = 24)
> fore.diff = as.numeric(fore1@forecast$seriesFor)
> fore.sigma = as.numeric(fore1@forecast$sigmaFor)
> ts.predict = temp[length(temp)] + cumsum(fore.diff)
> ts.predict = ts.predict + myTS.additive$figure
> ts.sigma = sqrt(cumsum(fore.sigma^2))
> tsup.sigma = ts.predict + ts.sigma
> tdown.sigma = ts.predict - ts.sigma
> plot(1:24, ts.predict, ylim=c(0,1.5), type = 'l', col = 'blue',
      xlab = "months", ylab = "temperature predict")
> lines(1:24, tsup.sigma, type = 'l', col = 'red')
> lines(1:24, tdown.sigma, type = 'l', col = 'red')
```


2.2 Figures

2.2.1 Data adjustment

Figure 1: HadCRUT4 Data Global Mean Time Series





Figure 2: Plot myTS.adjusted



2.2.2 ARIMA Model

Figure 3: Selecting arma order: acf of residuals of arma(2, 0, 1) & (2, 1, 4)




Figure 4: Selecting arma order: pacf of residuals of arma(2, 0, 1) & (2, 1, 4)




Figure 5: Forecast next 20 years(240 months) using arima(2, 1, 4)





Figure 6: Forecast within the sample compared with the actual data using arima(2, 1, 4)



2.2.3 ARFIMA Model

Figure 7: Plot arima2's residuals





Figure 8: acf of residuals of ARFIMA



2.2.4 More attempts on lag 24

Original series' acf and pacf

Figure 9: acf




Figure 10: pacf




Figure 11: res2's acf




Figure 12: res3's acf






Figure 13: final residuals: res4's acf



The failure of trying pure seasonal adjustment

Figure 14: res5's acf



2.2.5 GARCH Model

Figure 15: z




Figure 16: simulation of using `rnorm`




Figure 17: $\text{acf}(z)$




Figure 18: Empirical Density of Standardized Residuals




Figure 19: qqplot of z





Figure 20: forecast for sGARCH Model



2.3 Others

2.3.1 Details about *decompose()* used in seasonal adjustment

Type 'decompose' in R console, we can see the source code of this function. The process of 'type = additive' is listed below:

- the argument passed into `decompose()` is a 'ts' object
- denote the argument `ts(x, frequency = f)`. Create a filter using: `filter = c(0.5, rep(1, f - 1), 0.5)/f`.
- `trend = filter(x, filter)`
- `season = x - trend`, then compute `f` means of `season` with interval length `f`, the `f` means denoted by `figure`. Adjusting `figure` `figure = figure - mean(figure)`
- `seasonal` is just `length(x)/f` times repetition of `figure`.
- `random = x - seasonal - trend`