

ARMA-GARCH

SHIHENG SHEN

May 29, 2017

1 Data

1.1 Data Source


HadCRUT4 is a gridded dataset of global historical surface temperature anomalies relative to a 1961-1990 reference period. Data are available for each month since January 1850, on a 5 degree grid. The dataset is a collaborative product of the Met Office Hadley Centre and the Climatic Research Unit at the University of East Anglia.

`url = http://www.metoffice.gov.uk/hadobs/hadcrut4/data/current/time_series/HadCRUT.4.5.0.0.monthly_ns_avg.txt.`

1.2 Brief Data Description

The gridded data are a blend of the CRUTEM4 land-surface air temperature dataset and the HadSST3 sea-surface temperature (SST) dataset. The dataset is presented as an ensemble of 100 dataset realisations that sample the distribution of uncertainty in the global temperature record given current understanding of non-climatic factors affecting near-surface temperature observations. This ensemble approach allows characterisation of spatially and temporally correlated uncertainty structure in the gridded data, for example arising from uncertainties in methods used to account for changes in SST measurement practices, homogenisation of land station records and the potential impacts of urbanisation. The HadCRUT4 data are neither interpolated nor variance adjusted.

Below is a graph provided by www.metoffice.gov.uk:



1.3 Loading the Data into R


For our purpose, we are not interested in the spatial distribution, so we only use the *Global Mean*.

Loading the data into R:

```
library(curl)
tmpf <- tempfile()
curl_download(url, tmpf)
gtemp <- read.table(tmpf)[, 1:2]
temp = gtemp$V2[1:2004]
```

We only use the first two columns: *time and monthly global mean*. We pick the first 2004 observations, which is monthly data from 1850 ~ 2016, so that the data has complete periods.

The time series of *temp* is shown below:



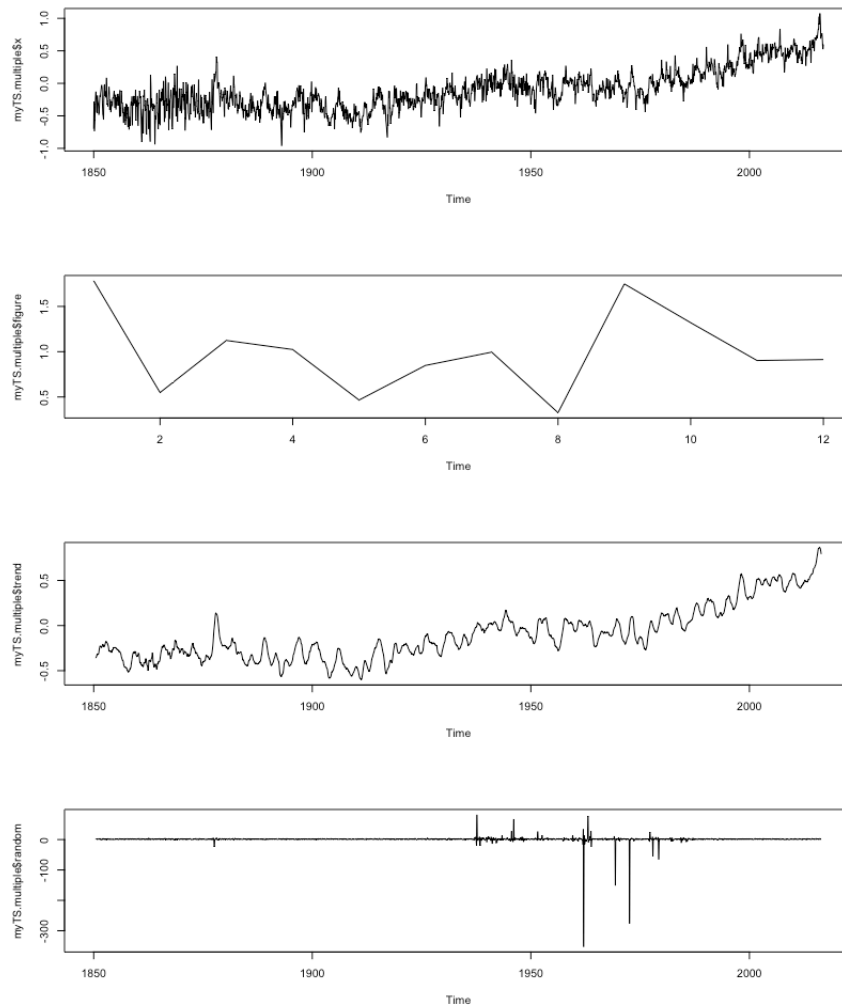
2 Dealing with the time series

2.1 Decompose

The first thing we might want to do is to wipe out the seasonal component. There're many approaches to do it in R. We choose to use the *decompose()* provided by *{stats}*. Looking at the time series, we think a additive model should be appropriate.

We can quickly checked that multiplicative model is not appropriate:

Figure 1: multiplicative model: original ts, seasonal component, time trend, residuals



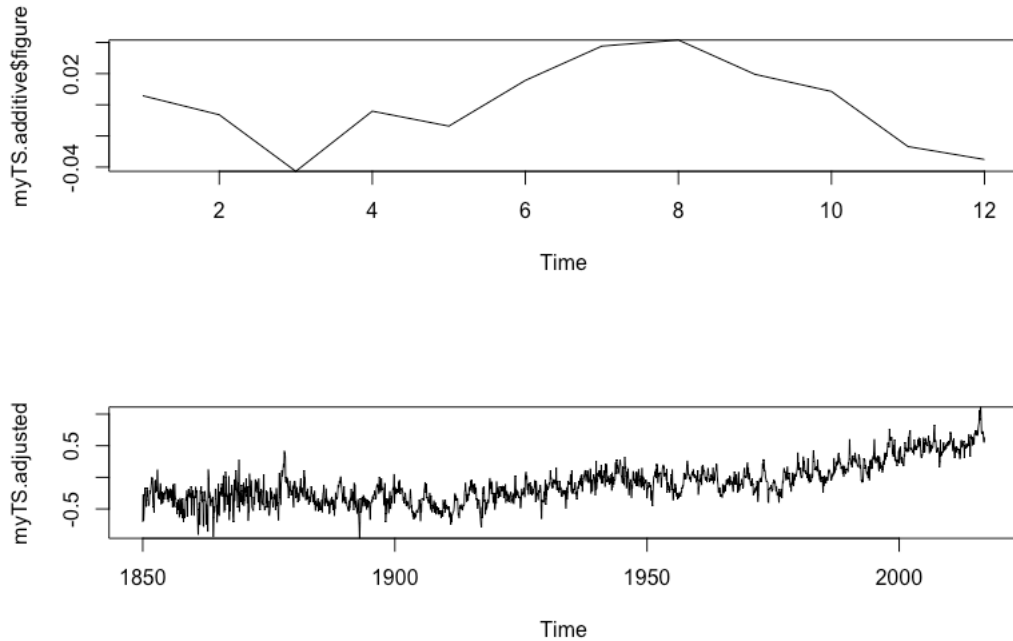
Obviously the random term is not easy to be modelled.

Therefore, decompose the time series with additive method in the following code:

```
library(TSA)
myTS = ts(as.numeric(temp), start = c(1850, 1), frequency = 12)
myTS.additive = decompose(myTS)
myTS.adjusted = myTS.additive$x - myTS.additive$seasonal
```

The *myTS.adjusted* now is free of seasonal components. The following graph illustrate the components.

Figure 2: additive model: seasonal component, adjusted series



Obviously, the time series is not stationary. For simplicity, use *diff()*.

```
dtemp = diff(myTS.adjusted)
```

```
> adf.test(dtemp)
```

Augmented Dickey-Fuller Test

```
data: dtemp
```

```
Dickey-Fuller = -16.175, Lag order = 12, p-value = 0.01
```

```
alternative hypothesis: stationary
```

So we can convince ourself that the time series *dtemp* is stationary.

2.2 Build a $ARMA(p, q)$ for *res*

For simplicity, we use *auto.arima()* from *{forecast}* to judge the order.

```
library(forecast)
```

```
> auto.arima(dtemp> auto.arima(dtemp)
```

```
Series: dtemp
```

```
ARIMA(2,0,4)(2,0,1)[12] with non-zero mean
```

```
Coefficients:
```

```
ar1      ar2      ma1      ma2      ma3      ma4      sar1      sar2      sma1      mean
```

| | | | | | | | | | | |
|------|--------|--------|---------|---------|--------|--------|--------|--------|---------|-------|
| | 0.5041 | 0.3585 | -1.0458 | -0.1600 | 0.1362 | 0.0780 | 0.7619 | 0.0779 | -0.7884 | 5e-04 |
| s.e. | 0.2431 | 0.2142 | 0.2423 | 0.3481 | 0.1082 | 0.0257 | 0.0926 | 0.0248 | 0.0910 | 2e-04 |

sigma^2 estimated as 0.01428: log likelihood=1417.03

AIC=-2812.05 AICc=-2811.92 BIC=-2750.43

first find the possible orders might be arma(2, 1) or arma(2, 4)

arma21.dtemp = arima(dtemp, c(2, 0, 1))

arma24.dtemp = arima(dtemp, c(2, 0, 4))

next examine the two models' residuals

> auto.arima(arma21.dtemp\$residuals)

Series: arma21.dtemp\$residuals

ARIMA(2,0,3)(2,0,1)[12] with zero mean

Coefficients:

| | | | | | | | | |
|------|--------|---------|---------|--------|--------|--------|--------|---------|
| | ar1 | ar2 | ma1 | ma2 | ma3 | sar1 | sar2 | sma1 |
| | 1.5859 | -0.6762 | -1.6089 | 0.6706 | 0.0506 | 0.7454 | 0.0748 | -0.7785 |
| s.e. | 0.1156 | 0.1065 | 0.1169 | 0.1264 | 0.0297 | 0.1209 | 0.0258 | 0.1181 |

sigma^2 estimated as 0.01429: log likelihood=1415.92

AIC=-2813.83 AICc=-2813.74 BIC=-2763.41

> auto.arima(arma24.dtemp\$residuals)

Series: arma24.dtemp\$residuals

ARIMA(0,0,0) with zero mean

sigma^2 estimated as 0.01435: log likelihood=1408.31

AIC=-2814.63 AICc=-2814.62 BIC=-2809.02

Therefore it should be appropriate to choose *arma(2, 4)* for series *dtemp*.

my.arma = arma24.dtemp

> my.arma

Call:

arima(x = dtemp, order = c(2, 0, 4))

Coefficients:

| | | | | | | | |
|------|--------|--------|---------|---------|--------|--------|-----------|
| | ar1 | ar2 | ma1 | ma2 | ma3 | ma4 | intercept |
| | 0.5204 | 0.3247 | -1.0536 | -0.1269 | 0.1168 | 0.0755 | 5e-04 |
| s.e. | 0.2733 | 0.2356 | 0.2727 | 0.3828 | 0.1117 | 0.0265 | 2e-04 |

sigma^2 estimated as 0.01435: log likelihood = 1407.67, aic = -2801.34

```
res = my.arma$residuals

> Box.test(res, type = 'Ljung-Box')

Box-Ljung test


data:  res
X-squared = 0.0058853, df = 1, p-value = 0.9388
```

The $ARMA(2,4)$ success in making the arma residual independent.

2.3 GARCH

One can be pretty satisfied with the results right now. Yet Looking at the plot of arma residual:

Figure 3: arma residual



It makes us wondering if it has $ARCH$ effect. Using *arch.test* from *{aTSA}* to perform the $LM - Test$:

```
library(aTSA)
```

```
> arch.test(my.arma)
ARCH heteroscedasticity test for residuals
alternative: heteroscedastic
```

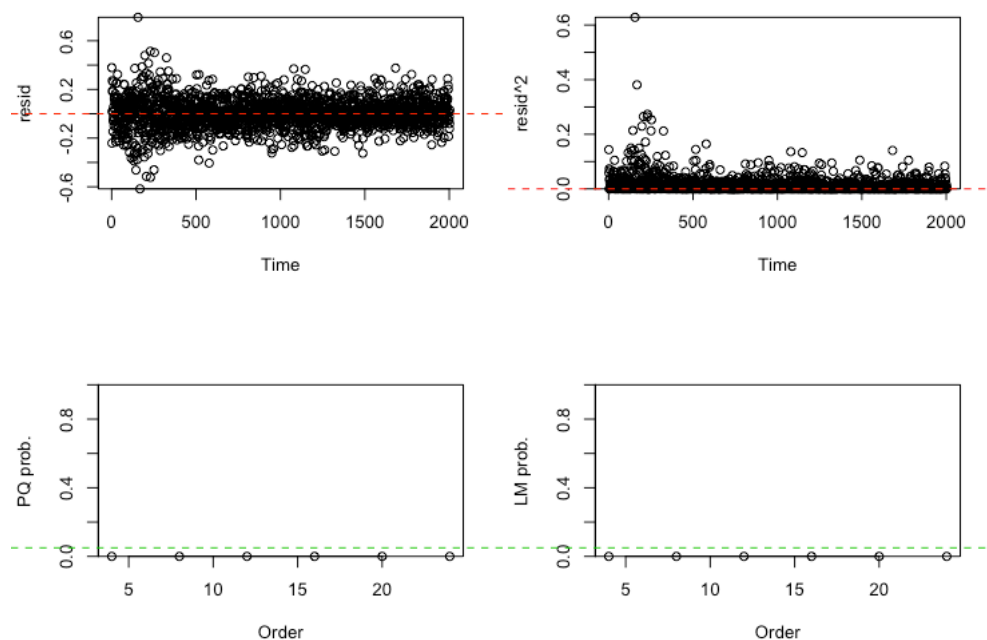
Portmanteau-Q test:

| | order | PQ | p.value |
|------|-------|-------|---------|
| [1,] | 4 | 99.4 | 0 |
| [2,] | 8 | 116.9 | 0 |
| [3,] | 12 | 411.5 | 0 |
| [4,] | 16 | 475.6 | 0 |
| [5,] | 20 | 495.9 | 0 |
| [6,] | 24 | 755.8 | 0 |

Lagrange-Multiplier test:

| | order | LM | p.value |
|------|-------|------|---------|
| [1,] | 4 | 1418 | 0 |
| [2,] | 8 | 697 | 0 |
| [3,] | 12 | 418 | 0 |
| [4,] | 16 | 212 | 0 |
| [5,] | 20 | 168 | 0 |
| [6,] | 24 | 137 | 0 |

Figure 4: arma residual



It supports our doubt. Therefore we have the motive to build a *GARCH* model. We decide to use the package *rugarch*.

$$\sigma_t^2 = (w + \sum_{j=1}^m \zeta_j v_{jt}) + \sum_{j=1}^q \alpha_j \varepsilon_{t-j}^2 + \sum_{j=1}^p \beta_j \sigma_{t-j}^2$$

So we wrote a function:

```
library(rugarch)
my_sGARCH_test <- function(p, q, m, n, ts.data = res)
{
# I use include.mean = FALSE after trying TRUE
# to find out insignificance
  myspec=ugarchspec(variance.model = list(model = "sGARCH", garchOrder = c(p, q)),
    mean.model = list(armaOrder = c(m, n), include.mean = FALSE),
    distribution.model = "norm")
  myfit=ugarchfit(myspec,data=ts.data, solver="solnp")
  return(myfit)
}
```

After trying a few times from (1,0) to (5,5), *GARCH(1,1)*, *ARIMA(2, 0, 4)* is the most satisfying model.

```
fit1 = my_sGARCH_test(1, 1, 2, 4, dtemp)
```

```
> fit1
```

```
*-----*
*          GARCH Model Fit          *
*-----*
```

Conditional Variance Dynamics

```
-----
GARCH Model : sGARCH(1,1)
Mean Model  : ARFIMA(2,0,4)
Distribution : norm
```

Optimal Parameters

```
-----
      Estimate Std. Error   t value Pr(>|t|)
ar1    -0.089169   0.038250   -2.33121 0.019742
ar2     0.763456   0.026364   28.95834 0.000000
ma1    -0.390348   0.019274  -20.25223 0.000000
ma2    -0.880951   0.000338 -2603.65163 0.000000
ma3     0.297715   0.026635   11.17774 0.000000
ma4     0.012983   0.050864    0.25525 0.798527
omega   0.000082   0.000034    2.40136 0.016334
alpha1  0.023019   0.003610    6.37686 0.000000
beta1   0.970582   0.004693   206.83542 0.000000
```


Robust Standard Errors:

| | Estimate | Std. Error | t value | Pr(> t) |
|--------|-----------|------------|-------------|----------|
| ar1 | -0.089169 | 0.083587 | -1.06678 | 0.286070 |
| ar2 | 0.763456 | 0.055864 | 13.66623 | 0.000000 |
| ma1 | -0.390348 | 0.041125 | -9.49179 | 0.000000 |
| ma2 | -0.880951 | 0.000796 | -1106.83968 | 0.000000 |
| ma3 | 0.297715 | 0.059384 | 5.01339 | 0.000001 |
| ma4 | 0.012983 | 0.117876 | 0.11014 | 0.912296 |
| omega | 0.000082 | 0.000036 | 2.28917 | 0.022069 |
| alpha1 | 0.023019 | 0.003647 | 6.31184 | 0.000000 |
| beta1 | 0.970582 | 0.003730 | 260.21194 | 0.000000 |

LogLikelihood : 1512.547

Information Criteria

| | |
|--------------|---------|
| Akaike | -1.5013 |
| Bayes | -1.4761 |
| Shibata | -1.5013 |
| Hannan-Quinn | -1.4921 |

Weighted Ljung-Box Test on Standardized Residuals

| | statistic | p-value |
|----------------------------|-----------|---------|
| Lag[1] | 0.5981 | 0.4393 |
| Lag[2*(p+q)+(p+q)-1][17] | 7.2782 | 0.9990 |
| Lag[4*(p+q)+(p+q)-1][29] | 18.4551 | 0.1194 |
| d.o.f=6 | | |
| H0 : No serial correlation | | |

Weighted Ljung-Box Test on Standardized Squared Residuals

| | statistic | p-value |
|-------------------------|-----------|-----------|
| Lag[1] | 31.45 | 2.045e-08 |
| Lag[2*(p+q)+(p+q)-1][5] | 38.17 | 1.533e-10 |
| Lag[4*(p+q)+(p+q)-1][9] | 51.67 | 3.408e-13 |
| d.o.f=2 | | |

Weighted ARCH LM Tests

| | Statistic | Shape | Scale | P-Value |
|-------------|-----------|-------|-------|-----------|
| ARCH Lag[3] | 0.009122 | 0.500 | 2.000 | 9.239e-01 |
| ARCH Lag[5] | 11.200028 | 1.440 | 1.667 | 3.341e-03 |
| ARCH Lag[7] | 20.729545 | 2.315 | 1.543 | 4.158e-05 |

Nyblom stability test

Joint Statistic: 2.8481

Individual Statistics:

ar1 0.24868
ar2 0.59898
ma1 0.17360
ma2 0.21711
ma3 0.06382
ma4 0.10027
omega 0.09915
alpha1 0.41690
beta1 0.19806

Asymptotic Critical Values (10% 5% 1%)

Joint Statistic: 2.1 2.32 2.82

Individual Statistic: 0.35 0.47 0.75

Sign Bias Test

| | t-value | prob | sig |
|--------------------|---------|-----------|-----|
| Sign Bias | 0.405 | 6.856e-01 | |
| Negative Sign Bias | 3.660 | 2.586e-04 | *** |
| Positive Sign Bias | 4.500 | 7.193e-06 | *** |
| Joint Effect | 33.651 | 2.347e-07 | *** |

Adjusted Pearson Goodness-of-Fit Test:


| group | statistic | p-value(g-1) |
|-------|-----------|-----------------|
| 1 | 20 | 53.37 4.122e-05 |
| 2 | 30 | 66.07 1.025e-04 |
| 3 | 40 | 90.72 5.294e-06 |
| 4 | 50 | 98.27 3.776e-05 |

Elapsed time : 0.3975561

Substract the standardized(w.r.t. the variance model) residuals z , which is $z = \frac{residuals(fit)}{sigma(fit)}$.

`z = residuals(fit1) / sigma(fit1)`
`plot.ts(z)`

Figure 5: z




We can compute the mean and variance of z :

```
> mean(z)
[1] 0.0289858
> var(z)
      [,1]
[1,] 1.012043
> length(z)
[1] 2003
> plot(rnorm(2003, 0.0289858, 1.012043))
```

And then, just for fun, plot a normal sample series with the same parameters:

Figure 6: simulation of using rnorm



At least a human can't distinguish between them anymore. But what we most care about is whether standardized squared residuals can pass the *LM – Test*(subtracting from the above long result):

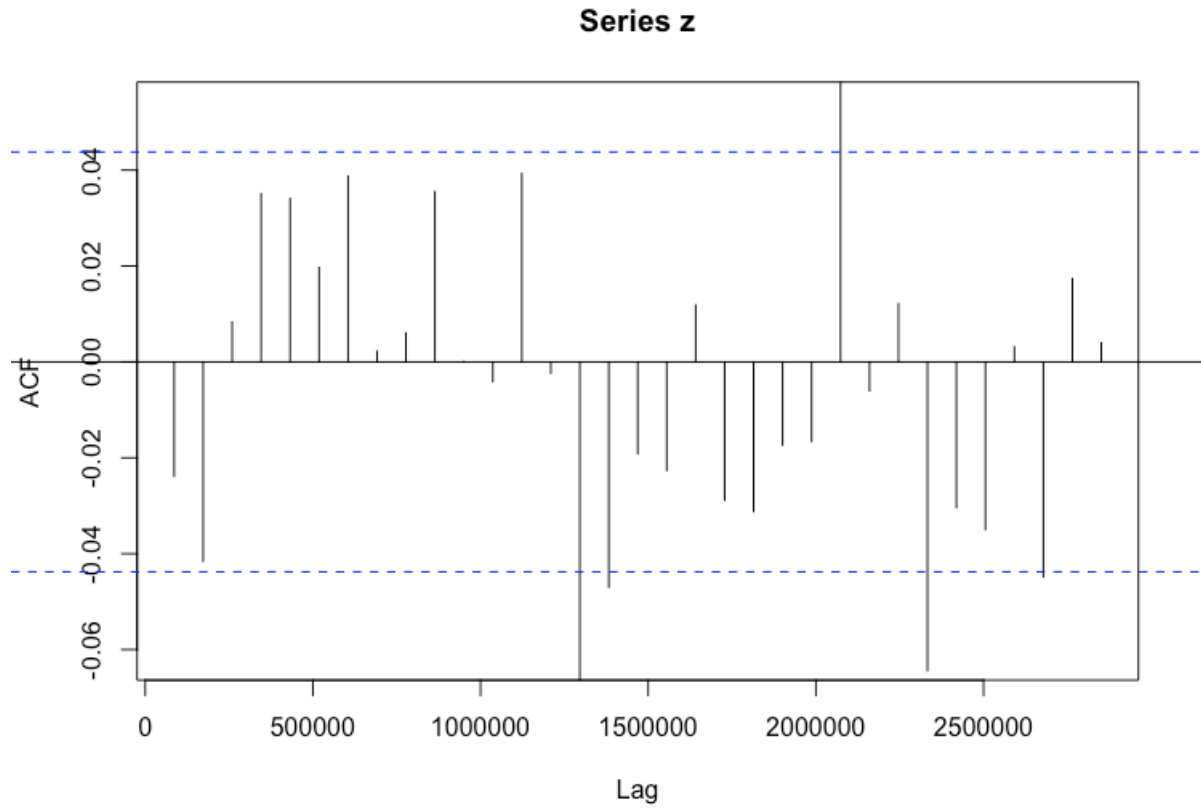
Weighted ARCH LM Tests

| | Statistic | Shape | Scale | P-Value |
|-------------|-----------|-------|-------|-----------|
| ARCH Lag[3] | 0.009122 | 0.500 | 2.000 | 9.239e-01 |
| ARCH Lag[5] | 11.200028 | 1.440 | 1.667 | 3.341e-03 |
| ARCH Lag[7] | 20.729545 | 2.315 | 1.543 | 4.158e-05 |

We can see our model successfully wipes out ARCH effect at small lags, but fails at larger lags. This is the best 'sGARCH' can give with respect to our data. Maybe we should try long memory models instead of *diff()* at first.


Graphical Diagnostics:

Figure 7: $\text{acf}(z)$




This is a prove that we should have used a long memory model.
rugarch provide a function to plot a "uGARCHfit" oboject.
ex. density smmothed to normal distribution:

Figure 8: Empirical Density of Standardized Residuals



qqplot:

Figure 9: qqplot of z




We can see that z is skewed, which might be correlated with long lags' heterodasticity.

Forecast

```
fore1 = ugarchforecast(fit1, n.ahead = 24)
fore.diff = as.numeric(fore1@forecast$seriesFor)
fore.sigma = as.numeric(fore1@forecast$sigmaFor)
ts.predict = temp[length(temp)] + cumsum(fore.diff)
ts.predict = ts.predict + myTS.additive$figure
ts.sigma = sqrt(cumsum(fore.sigma^2))
tsup.sigma = ts.predict + ts.sigma
tsdown.sigma = ts.predict - ts.sigma
plot(1:24, ts.predict, ylim=c(0,1.5), type = 'l', col = 'blue', xlab = "months", ylab = "ter")
lines(1:24, tsup.sigma, type = 'l', col = 'red')
lines(1:24, tsdown.sigma, type = 'l', col = 'red')
```

Figure 10: qqplot of z



Conclusion: 'sGARCH' have improved the model in the intuitive sense and to some extent fix the heterodasticity.